

COMP353 Databases

**Relational Algebra (RA)
for Relational Data Model**

Relational Algebra (RA)

- *Database Query languages* are specialized languages to ask for information (**queries**) in DB.
- **Relational Algebra (RA)** is a query language associated with the relational data model.
- Queries in RA are expressions using a collection of operators on relations in the DB.
- The input(s) and output of a RA query are relations
- A query is ***evaluated*** using the ***current instance*** of the input relations to produce the output

Operations in “standard” RA

- The well-known **set operations**
 - ✓ **Union (\cup)**
 - ✓ **Intersection (\cap)**
 - ✓ **Difference ($-$)**
- **Special DB operations** that select “parts” of a relation instance
 - **Selection (σ)** – selects some rows (tuples) & discards the rest
 - **Projection (π)** – selects some columns (attributes) & discards the rest
- Operations that “combine” the tuples from the argument relations
 - ✓ **Cartesian product (\times)** – pairs the tuples in all possible ways
 - **Join (\bowtie)** – pairs particular tuples from the two input relations
- A unary operation to **rename** relations, called **Rename (ρ)**

Note: The output of a RA expression is an “unnamed” relation/set, i.e.,
RA expressions return **sets**, whereas **SQL** returns **multisets (bags)**

Compatibility Requirement

- We can apply the **set operators** of **union**, **intersection**, and **difference** to instances of relations **R** and **S** if **R** and **S** are *compatible*, that is, they have “the same” schemas.
- **Definition:** Relations **S**(**A**₁,...,**A**_n) and **R**(**B**₁,...,**B**_m) are compatible if:
 - (1) **n=m** and
 - (2) **type(A_i) = type(B_i)** (or compatible types), for all **1 ≤ i ≤ n**.

Set Operations on Relations

Let **R** and **S** be relation schemas, and **r** and **s** be any instances of them.

- The **union** of **r** and **s** is the set of all tuples that appear in either one or both. Each tuple **t** appears only once in the union, even if it appears in both; $\mathbf{r} \cup \mathbf{s} = \{t \mid t \in \mathbf{r} \vee t \in \mathbf{s}\}$
- The **intersection** of **r** and **s**, is the set of all tuples that appear in both; $\mathbf{r} \cap \mathbf{s} = \{t \mid t \in \mathbf{r} \wedge t \in \mathbf{s}\}$
- The **difference** of **r** and **s**, is the set of all tuples that appear in **r** but not in **s**; $\mathbf{r} - \mathbf{s} = \{t \mid t \in \mathbf{r} \wedge t \notin \mathbf{s}\}$
- **Commutative** operations; $\mathbf{r} \mathbf{Op} \mathbf{s} = \mathbf{s} \mathbf{Op} \mathbf{r}$

Note: Set difference (−) is not commutative, i.e., $(\mathbf{r} - \mathbf{s}) \neq (\mathbf{s} - \mathbf{r})$

Example

Relation Schema: **Star** (name, address, gender, birthdate)

**Instance r
of Star:**

Name	Address	Gender	Birthdate
Carrie Fisher	123 Maple	F	9/9/99
Mark Hamill	456 Oak rd.	M	8/8/88

**Instance s
of Star:**

Name	Address	Gender	Birthdate
Carrie Fisher	123 Maple	F	9/9/99
Harrison Ford	789 Palm rd.	M	7/7/77

r ∪ s:

Name	Address	Gender	Birthdate
Carrie Fisher	123 Maple	F	9/9/99
Mark Hamill	456 Oak rd.	M	8/8/88
Harrison Ford	789 Palm rd.	M	7/7/77

Example

Relation Schema: **Star** (name, address, gender, birthdate)

**Instance r
of Star:**

Name	Address	Gender	Birthdate
Carrie Fisher	123 Maple	F	9/9/99
Mark Hamill	456 Oak rd.	M	8/8/88

**Instance s
of Star:**

Name	Address	Gender	Birthdate
Carrie Fisher	123 Maple	F	9/9/99
Harrison Ford	789 Palm rd.	M	7/7/77

r \cap s:

Name	Address	Gender	Birthdate
Carrie Fisher	123 Maple	F	9/9/99

Example

Relation Schema: **Star** (name, address, gender, birthdate)

**Instance r
of Star:**

Name	Address	Gender	Birthdate
Carrie Fisher	123 Maple	F	9/9/99
Mark Hamill	456 Oak rd.	M	8/8/88

**Instance S
of Star:**

Name	Address	Gender	Birthdate
Carrie Fisher	123 Maple	F	9/9/99
Harrison Ford	789 Palm rd.	M	7/7/77

r – S:

Name	Address	Gender	Birthdate
Mark Hamill	456 Oak rd.	M	8/8/88

Example

Relation Schema: **Star** (name, address, gender, birthdate)

**Instance r
of Star:**

Name	Address	Gender	Birthdate
Carrie Fisher	123 Maple	F	9/9/99
Mark Hamill	456 Oak rd.	M	8/8/88

**Instance s
of Star:**

Name	Address	Gender	Birthdate
Carrie Fisher	123 Maple	F	9/9/99
Harrison Ford	789 Palm rd.	M	7/7/77

s – r:

Name	Address	Gender	Birthdate
Harrison Ford	789 Palm rd.	M	7/7/77

Projection (π)

- Let \mathbf{R} be a relation schema.
- The **projection** operation (π) is used to produce, from any instance \mathbf{r} of \mathbf{R} , a new relation that includes listed “columns” of \mathbf{R}
- The output of $\pi_{A_1, A_2, \dots, A_j}(\mathbf{r})$ is a relation with columns A_1, A_2, \dots, A_j , in this order.
- Note: The subscript of π is a *list*, which defines the structure of the output as the ordered tuple (A_1, A_2, \dots, A_j) .

Example

Relation Schema: **Movie**(title, year, length, filmType, studioName, producer)

Instance
movie
Of Movie:

title	year	length	filmType	studioName	producer
Star wars	1977	124	color	Fox	12345
Mighty Ducks	1991	104	color	Disney	67890
Wayne's World	1992	95	color	Paramount	99999

Query: π title, year, length(movie)

title	year	length
Star wars	1977	124
Mighty Ducks	1991	104
Wayne's World	1992	95

Example

Relation Schema: **Movie**(title, year, length, filmType, studioName, producer)

Instance
movie
Of Movie:

title	year	length	filmType	studioName	producer
Star wars	1977	124	color	Fox	12345
Mighty Ducks	1991	104	color	Disney	67890
Wayne's World	1992	95	color	Paramount	99999

Query: π filmType(movie)

Result:

filmType
color

Selection (σ)

- The **selection** operator (σ), applied to an instance **r** of relation **R**, returns a subset of **r**
- We denote this operation/query by $\sigma_c(r)$
- The output includes tuples satisfying condition **C**
- The schema of the output is the same as **R**

Example

Relation Schema: **Movie**(title, year, length, filmType, studioName, producer)

**Instance movie
of Movie:**

title	year	length	filmType	studioName	producer
Star wars	1977	124	color	Fox	12345
Mighty Ducks	1991	104	color	Disney	67890
Wayne's World	1992	95	color	Paramount	99999

Query: $\sigma_{\text{length} \geq 100}(\text{movie})$

Result:

title	year	length	filmType	studioName	producer
Star wars	1977	124	color	Fox	12345
Mighty Ducks	1991	104	color	Disney	67890

Example

Relation: Movie(title, year, length, filmType, studioName, producer)

**Instance
movie
of Movie:**

title	year	length	filmType	studioName	producer
Star wars	1977	124	color	Fox	12345
Mighty Ducks	1991	104	color	Disney	67890
Wayne's World	1992	95	color	Paramount	99999

Query: $\sigma_{\text{length} \geq 100 \text{ AND studioName} = \text{'Fox'}}(\text{movie})$

Result:

title	year	length	filmType	studioName	producer
Star wars	1977	124	color	Fox	12345

Cartesian Product (\times)

- Let **R** and **S** be relation schemas, and **r** and **s** be any instances of **R** and **S**, respectively.
- The **Cartesian Product** of **r** and **s** is the set of all tuples obtained by “concatenating” the tuples in **r** and **s**. Formally, $\mathbf{r} \times \mathbf{s} = \{ t_1.t_2 \mid t_1 \in r \wedge t_2 \in s \}$
- The schema of result is the “union” of **R** and **S**
 - If **R** and **S** have some attributes in common, we need to invent new names for identical names, e.g., use **R.B** and **S.B**, if **B** appears in both **R** and **S**

Example

Instance r of R:

A	B
1	2
3	4

Instance s of S:

B	C	D
2	5	6
4	7	8
9	10	11

$r \times s$:

A	R.B	S.B	C	D
1	2	2	5	6
1	2	4	7	8
1	2	9	10	11
3	4	2	5	6
3	4	4	7	8
3	4	9	10	11

Theta-join (θ)

- Suppose \mathbf{R} and \mathbf{S} are relation schemas, \mathbf{r} is an instance of \mathbf{R} , and \mathbf{s} is an instance of \mathbf{S} . The **theta-join** of \mathbf{r} and \mathbf{s} is the set of all tuples obtained from concatenating all $t_1 \in \mathbf{r}$ and $t_2 \in \mathbf{s}$, such that t_1 and t_2 *satisfy some condition \mathbf{C}*
- We denote θ -join by $\mathbf{r} \triangleright\triangleleft_{\mathbf{C}} \mathbf{s}$
- The schema of the result is the same as the schema of $\mathbf{R} \times \mathbf{S}$ (i.e., the union of \mathbf{R} and \mathbf{S})
- \mathbf{C} is a Boolean expression, simple or complex, as in operation σ

Example

Instance r of R:

A	B	C
1	2	3
6	5	8
9	7	11

Instance s of S:

B	C	D
2	3	4
2	3	5
7	8	10

$r \bowtie_{A < D} s$:

A	R.B	R.C	S.B	S.C	D
1	2	3	2	3	4
1	2	3	2	3	5
1	2	3	7	8	10
6	5	8	7	8	10
9	7	11	7	8	10

Equi-join

- The **equi-join** operator, is a special case of θ -join, in which we may only use the equality relation (=) in condition **C**
- It is denoted as $\mathbf{r} \bowtie_{\mathbf{c}} \mathbf{s}$ (i.e., the same as θ -join)
- The schema of the output is the same as that of θ -join

Example

Instance r of R:

A	B	C
1	2	3
6	5	8
9	7	11

Instance s of S:

B	C	D
2	3	4
2	3	5
7	8	10

$r \bowtie_{R.C = S.C} s$

A	R.B	R.C	S.B	S.C	D
1	2	3	2	3	4
1	2	3	2	3	5
6	5	8	7	8	10

Natural Join (\bowtie)

- **Natural join**, is a special case of equi-join, where the equalities are not explicitly specified, rather they are assumed implicitly on the common attributes of **R** and **S**
- We denote this natural join operation by $\mathbf{r} \bowtie \mathbf{s}$
- The schema of the output is similar to that of equi-join, except that each common attribute appears only once.

Note: If **R** and **S** do not have any common attribute, then the join operation becomes Cartesian product.

Example

Instance r of R:

A	B	C
1	2	3
6	2	8
9	7	3

Instance s of S:

B	C	D
2	3	4
2	3	5
7	8	10

$r \bowtie s$:

A	B	C	D
1	2	3	4
1	2	3	5

Expressing Queries in RA

- Every standard RA operation has relation(s) as argument(s) and produces a relation (set) as the output
(Exception is the sort operator τ)
- This property of RA operations (that inputs and outputs are relations) makes it possible to formulate/express any query by *composing/nesting/grouping* subqueries.
- We can use parentheses for *grouping*, in order to improve *clarity* and *readability*

Example: RA Query

- Relation schema:

Movie (title, year, length, filmType, studioName)

- Query: List the title and year of every movie made by Fox studio whose length is at least 100 minutes?

- One way to express this query in RA is:


$\pi_{\text{title, year}} (\sigma_{\text{studioName} = \text{'Fox'} \text{ and } \text{length} \geq 100} (\text{movie}))$

- Another way:

- Select those **movie** tuples that have *length* ≥ 100
- Select those **movie** tuples that have *studioName* = 'Fox'
- Find the **intersection** of the above two results
- Then project on the attributes **title** and **year**

Example: RA Query

$\pi_{title, year} (\sigma_{studioName = 'Fox' \text{ and } length \geq 100} (movie))$



$\pi_{title, year} (\sigma_{length \geq 100} (movie) \cap \sigma_{studioName = 'Fox'} (movie))$

Example: RA Query

- Relation schema:

Movie (title, year, length, filmType, studioName)

StarsIn (title, year, starName)

- Query: List the **names** of the stars of movies of length ≥ 100 minutes long.
- One expression in RA for this query:
 - Select **movie** tuples of **length** ≥ 100
 - Join the result with relation **StarsIn**
 - Project on the attribute **starName**
- **Exp1:** $\pi_{starName}(\sigma_{length \geq 100}(\text{movie}) \bowtie \text{starsIn})$
- **Another solution:** $\pi_{starName}(\sigma_{length \geq 100}(\text{movie} \bowtie \text{starsIn}))$

Renaming Operator (ρ)

- To **control** manipulating the names of the attributes in formulating queries in relational algebra, we may need **renaming** of relations. May do this for *convenience too*
- **The Renaming Operator** is denoted by $\rho_{s(A_1, A_2, \dots, A_n)}(r)$
- The result is a copy of the input relation instance **r**, but renamed to **s** and its attributes to A_1, \dots, A_n , in that order.
- Use $\rho_s(r)$ to give relation **r** a new name **s** (**with the same attributes in r**)

That is, in this case, schema of **s** is the same as that of **r**.

Example

- Query: $\pi_{starName}(\sigma_{length \geq 100}(\text{movie}) \bowtie \text{starsIn})$
- This query can be rewritten in 2 steps as follows:
 1. $\rho_M(\text{title, year, length, filmType, studioName})(\sigma_{length \geq 100}(\text{movie}))$
or even simpler as: $\rho_M(\sigma_{length \geq 100}(\text{movie}))$ if used in the same formula
 2. Or use $M := \sigma_{length \geq 100}(\text{movie})$ as a separate formula and then formulate the query as: $\pi_{starName}(M \bowtie \text{starsIn})$
- Consider $\text{takes}(\underline{\text{sid}}, \underline{\text{cid}}, \text{grade})$
- Query: Find ID of every student who has taken at least 2 courses.
- $\pi_{\text{takes.sid}}(\sigma_{(\text{takes.sid} = T.\text{sid}) \text{ and } (\text{takes.cid} \neq T.\text{cid})}(\text{takes} \times \rho_T(\text{takes})))$

Dependent and Independent Operations

- Some RA operations can be expressed based on other operations. Examples include:
 - $r \cap s = r - (r - s)$
 - $r \bowtie_c s = \sigma_c (r \times s)$
 - $r \bowtie s = \pi_L(\sigma_{r.A1 = s.A1 \text{ AND... AND } r.An = s.An} (r \times s)),$
where L is the list of attributes in R followed by those attributes in S that are not in R , and $A1, \dots, An$ are the common attributes of R and S

Relational Algebra with Bag Semantics

- Relations stored in DB are called **base** relations/tables.
- Base relations are normally sets; no duplicates.
- In some situations, e.g., during query processing, it is allowed for relations to have duplicate tuples.
- If duplicates are allowed in a collection, it is called **bag/multiset**.

Instance
r of R:

A	B	C
1	2	3
6	5	8
6	5	8
1	2	3
9	7	11

Here, r is a bag

Why Bags?

■ 1. Faster projection operations

- Bag projection is faster, since otherwise returning distinct values is expensive (as we need sorting for duplicate elimination).

Another example: *Computing the bag union $(\mathbf{r} \cup^{\mathbf{B}} \mathbf{s})$ is much cheaper than computing the standard set union $\mathbf{r} \cup \mathbf{s}$.*

Formally, if \mathbf{r} and \mathbf{s} have n and m tuples, then the bag and set union operations will cost $O(n+m)$ and $O(n*m)$, respectively.

■ 2. Correct computation with some aggregation

- For example, to compute the *average* of values for attribute A in the previous relation, we must consider the bag of those values

Set Operations on Bags

- $r \cup^B s$, the **bag union** of r and s , is the bag of tuples that are in r , in s , or in both. If a tuple t appears n times in r , and m times in s , then t appears $n+m$ times in bag $r \cup^B s$

$$r \cup^B s = \{ t:k \mid t:n \in r \wedge t:m \in s \wedge k = n+m \}$$

- $r \cap^B s$, the **bag intersection** of r and s , is the bag of tuples that appear in both r and s . If a tuple t appears n times in r , and m times in s , then the number of occurrences of t in bag $r \cap^B s$ is $\min(n,m)$

$$r \cap^B s = \{ t:k \mid t:n \in r \wedge t:m \in s \wedge k = \min(n,m) \}$$

- $r -^B s$, the **bag difference** of r and s is defined as follows:

$$r -^B s = \{ t:k \mid t:n \in r \wedge t:m \in s \wedge k = \max(0, n-m) \}$$

$$s -^B r = \{ t:k \mid t:n \in r \wedge t:m \in s \wedge k = \max(0, m-n) \}$$

Example

Bag r:

A	B
1	2
3	4
1	2
1	2

Bag s:

A	B
1	2
3	4
3	4
5	6

$r \cup^B s$:

A	B
1	2
3	4
1	2
1	2
1	2
3	4
3	4
5	6

Example

Bag r:

A	B
1	2
3	4
1	2
1	2

Bag s:

A	B
1	2
3	4
3	4
5	6

$r \cap^B s$:

A	B
1	2
3	4

Example

Bag r:

A	B
1	2
3	4
1	2
1	2

Bag s:

A	B
1	2
3	4
3	4
5	6

$r \text{ --}^B s$:

A	B
1	2
1	2

Example

Bag r:

A	B
1	2
3	4
1	2
1	2

Bag s:

A	B
1	2
3	4
3	4
5	6

$s \stackrel{B}{-} r$:

A	B
3	4
5	6

Bag Projection π^B

- Let \mathbf{R} be a relation scheme, and \mathbf{r} be a collection of tuples over \mathbf{R} , which could have duplicates.

The **bag projection** operator is used to produce, from \mathbf{r} , a bag of tuples over some of \mathbf{R} .

- Even when \mathbf{r} does not have duplicates, we may get duplicates when projecting on some attributes of \mathbf{R} .

That is, π^B does not eliminate the duplicates and hence corresponds exactly to the `SELECT` clause in SQL.

Example

Bag r:

A	B	C
1	2	5
3	4	6
1	2	7
1	2	8

$\pi_{A, B}^B(r)$:

A	B
1	2
3	4
1	2
1	2

Example

Relation Schema: **movie**(title, year, length, filmType, studioName, producer)

Instance:

title	year	length	filmType	studioName	producer
Star wars	1977	124	color	Fox	12345
Mighty Ducks	1991	104	color	Disney	67890
Wayne's World	1992	95	color	Paramount	99999

$\pi_{\text{filmType}}^{\text{B}}(\text{movie})$:

filmType
color
color
color

Selection on Bags

- The selection operator σ_c applied to an instance r of relation R will return a subset of r
 - The tuples returned are those that satisfy the specified condition C (which involves attributes of R)
 - Duplicates are **not** eliminated from the result of a bag-selection

Note: The selection operation σ in RA is different from the **SELECT** clause in SQL

Example

Bag r:

A	B	C
1	2	5
3	4	6
1	2	7
1	2	7

$\sigma_{C \geq 6}(r)$:

A	B	C
3	4	6
1	2	7
1	2	7

Cartesian Product of Bags

- The **Cartesian Product** of bags **r** and **s** is the bag of tuples that can be formed by concatenating pairs of tuples, the first of which comes from **r** and the second from **s**. In symbols, $\mathbf{r} \times \mathbf{s} = \{ t_1.t_2 \mid t_1 \in \mathbf{r} \wedge t_2 \in \mathbf{s} \}$
- Each tuple of one relation is paired with each tuple of the other, regardless of whether it is a duplicate or not
- If a tuple t_1 appears **m** times in a relation **r**, and a tuple t_2 appears **n** times in relation **s**, then tuple $t_1.t_2$ appears **m*n** times in their bag-product, $\mathbf{r} \times \mathbf{s}$

Example

Bag r:

A	B
1	2
1	2

Bag s:

B	C
2	3
4	5
4	5

$r \times s$:

A	R.B	S.B	C
1	2	2	3
1	2	2	3
1	2	4	5
1	2	4	5
1	2	4	5
1	2	4	5

Join of Bags

- The bag join is computed in the same way as the standard join operation
- Duplicates are not eliminated in a bag join operation

Bag r:

A	B
1	2
1	2

Bag s:

B	C
2	3
4	5

$r \bowtie s$:

A	B	C
1	2	3
1	2	3

Constraints on Relations

- RA offers a convenient way to express a wide variety of constraints, e.g., referential integrity and FD's.
- There are **two ways** to express constraints in RA
 1. If r is an expression in RA, then the constraint $r = \emptyset$ says:
“ r has no tuples, i.e., or r is empty”
 2. If r and s are RA expressions, then the constraint $r \subseteq s$ says:
“every tuple in (the result of) r is in (the result of) s ”

These constraints hold also when r and s are bags.

Constraints on Relations

- Note that these two types of constraints are not independent. Why?
 - The constraint $\mathbf{r} \subseteq \mathbf{s}$ could also be written as $\mathbf{r} - \mathbf{s} = \emptyset$

This follows from the definition of “ $-$ ”, because $\mathbf{r} \subseteq \mathbf{s}$ **iff** $\mathbf{r} - \mathbf{s} = \emptyset$, meaning that there is no tuple in \mathbf{r} which is not in \mathbf{s}

Referential Integrity Constraints

- Referential integrity in relational data model means:
 - if there is a value v in a tuple t in a relation r , then it is expected that v appears in a particular component (attribute) of some tuple s in relation s
- E.g., if tuple (s,c,g) is in table **takes**(*sid*,*cid*,*grade*), then there must be a **student** with *sid* = s and a **course** with *cid* = c such that s has taken c
- IOW, the mentions of values s and c in **takes** “refers” to some values outside this relation, and these values **must** exist*

Example

- Relation schemas:

Movie (title, year, length, filmType)

StarsIn (title, year, starName)

- Constraint:

the **title** and **year** of every movie that appears in relation **starsIn** *must* appear also in **movie**; otherwise there is a violation in referencing in **starsIn**

- Query in RA:

- $\pi_{\text{title, year}}(\mathbf{starsIn}) \subseteq \pi_{\text{title, year}}(\mathbf{movie})$

or equivalently

- $\pi_{\text{title, year}}(\mathbf{starsIn}) - \pi_{\text{title, year}}(\mathbf{movie}) = \emptyset$

Functional Dependencies

- **Any** functional dependency $X \rightarrow Y$ can be expressed as an expression in RA
- Example:
Consider the relation schema:
Star (name, address, gender, birthdate)
- How to express the FD: **name** \rightarrow **address** in RA?

Functional Dependencies

- Relation schema:
Star (name, address, birthdate)
- With the FD: **name** → **address**
- The **idea** is that if we construct all pairs of **star** tuples, we **must not** find a pair that agree on **name** but disagree on **address**
- To “construct” the pairs in RA, we use **Cartesian product**, and to find pairs that violate this FD, we use **selection**
- We are then ready to express this **FD** by equating the result to \emptyset , as follows...

Example

Star:

Name	Address	Birthdate
Carrie Fisher	123 Maple	9/9/99
Mark Hamill	456 Oak rd.	8/8/88
Harrison Ford	789 Palm rd.	7/7/77

$\rho_{S1}(\textit{name}, \textit{address}, \textit{birthdate})(\textbf{star})$

Name	Address	Birthdate
Carrie Fisher	123 Maple	9/9/99
Mark Hamill	456 Oak rd.	8/8/88
Harrison Ford	789 Palm rd.	7/7/77

$\rho_{S2}(\textit{name}, \textit{address}, \textit{birthdate})(\textbf{star})$

Name	Address	Birthdate
Carrie Fisher	123 Maple	9/9/99
Mark Hamill	456 Oak rd.	8/8/88
Harrison Ford	789 Palm rd.	7/7/77

Example

$s1 \times s2$:

S1.Name	S1.Address	S1.Birthdate	S2.Name	S2.Address	S2.Birthdate
Carrie Fisher	123 Maple	9/9/99	Carrie Fisher	123 Maple	9/9/99
Carrie Fisher	123 Maple	9/9/99	Mark Hamill	456 Oak rd.	8/8/88
Carrie Fisher	123 Maple	9/9/99	Harrison Ford	789 Palm rd.	7/7/77
Mark Hamill	456 Oak rd.	8/8/88	Carrie Fisher	123 Maple	9/9/99
Mark Hamill	456 Oak rd.	8/8/88	Mark Hamill	456 Oak rd.	8/8/88
Mark Hamill	456 Oak rd.	8/8/88	Harrison Ford	789 Palm rd.	7/7/77
Harrison Ford	789 Palm rd.	7/7/77	Carrie Fisher	123 Maple	9/9/99
Harrison Ford	789 Palm rd.	7/7/77	Mark Hamill	456 Oak rd.	8/8/88
Harrison Ford	789 Palm rd.	7/7/77	Harrison Ford	789 Palm rd.	7/7/77

$\sigma_{S1.name=S2.name \text{ AND } S1.address \neq S2.address}(s1 \times s2) = \emptyset$

Functional Dependencies

- Relation schema:
Star (name, address, birthdate)
- With the FD: **name** \rightarrow **address**
- In RA:

$$\sigma_{S1.name=S2.name \text{ AND } S1.address \neq S2.address}(\rho_{S1}(\text{star}) \times \rho_{S2}(\text{star})) = \emptyset$$

Domain Constraints

- Relation schema:
Star (name, address, gender, birthdate)
- How to express the following constraint?
Valid values for **gender** are 'F' and 'M'
- In RA:
 - $\sigma_{\text{gender} \neq \text{'F'} \text{ AND } \text{gender} \neq \text{'M'}}(\text{star}) = \emptyset$
 - This is an example of *domain constraints*

Domain Constraints

- Relation schema:
Employee (eid, name, address, salary)
- How to express the constraint:
Maximum employee salaries is \$150,000
- In RA:
 - $\sigma_{\text{salary} > 150000}(\text{employee}) = \emptyset$

"For All" Queries (1)

- Given the database schema:

Student(Sid, Sname, Addr)

Course(Cid, Cname, Credits)

Enrolled (Sid, Cid)

- Consider the query:

*"Find students enrolled in **all** the courses."*

- A first attempt (below) fails!

$\pi_{sid}(\text{Enrolled})$

- This RA query returns those students enrolled in some courses.
- So, how to correctly express "For All" types of queries?

"For All" Queries (2)

- A *solution strategy* would be to:
 - First find the list of “all” students (**all guys**), from which we then subtract those who have not taken at least a course (**bad guys**)
 - Then “**good guys**”, would be “all guys” from which we remove the “bad guys”, i.e.,
Answer (Good guys) = All guys – Bad guys

"For All" Queries (3)

- Set of all students that we should consider:

All Courses := $\pi_{Cid}(\text{Course})$

All Students := $\pi_{Sid}(\text{Student})$

- Steps to find **students not enrolled in all courses**

1. Create all possible “student-course” pairs:

All: Student-Course Pairs := $\pi_{Sid}(\text{Student}) \times \pi_{Cid}(\text{Course})$

2. Extract the “actual” student-course pairs from **Enrolled**

3. Using 1 & 2, we then find students not enrolled in all courses:

Bad : $\pi_{Sid}(\pi_{Sid}(\text{Student}) \times \pi_{Cid}(\text{Course}) - \text{Enrolled})$

- **Answer**: **All** - **Bad**

The Division Operation (\div)

- The previous query can be expressed in RA using the **division** operator \div
 - Divide `Enrolled` by $\pi_{cid}(\text{Course})$
that is, $\text{Enrolled} \div \pi_{cid}(\text{Course})$
 - Schema of the result is $\{\text{Sid}, \text{Cid}\} - \{\text{Cid}\}$
- $R \div S$ requires that the attributes of S to be a subset of the attributes of R .
 - The schema of the output would be $R - S$

Example: Enrolled (student, sport)

Find students enrolled in all sports {Hockey, Football}.

Enrolled (Student, sport)

Jim	Hockey
Joe	Football
Jim	Football
Sue	Hockey

Example: Enrolled(student, sport)

$$\pi_{student}(\text{Enrolled}) \times \pi_{sport}(\text{Enrolled}) - \text{Enrolled} =$$

Jim	Hockey
Jim	Football
Joe	Hockey
Joe	Football
Sue	Hockey
Sue	Football

—

Jim	Hockey
Joe	Football
Jim	Football
Sue	Hockey

=

Joe	Hockey
Sue	Football

$$\pi_{student}(\text{Enrolled}) -$$

$$\pi_{student}(\pi_{student}(\text{Enrolled}) \times \pi_{sport}(\text{Enrolled}) - \text{Enrolled})$$

All

Bad

Jim
Joe
Sue

—

Joe
Sue

= ?



Jim is the only student enrolled in all sports

Another Example

- $r \div s = \pi_{R-S}(r) - \pi_{R-S}(\pi_{R-S}(r) \times s - r)$
- Consider the following DB schema in a banking application:
 - Customer(cid, name)
 - Branch(bid, district)
 - Account(cid, bid)
- Query: "Find the names of those customers who have an account in *every* branch in the Westmount area"
- Solution?
 - $\pi_{name}(\text{Customer} \bowtie \text{Account} \bowtie (\sigma_{\text{district} = \text{"Westmount"}}(\text{Branch})))$?
- **No**, this query returns those customers who have an account at some branch in Westmount, but not necessarily at all such branches.

Database:

Customer(cid, name), Branch(bid, district), Account(cid, bid)

- We can use the division operator \div
 - Find all customer-branch pairs (**cid**, **bid**) for which customer (cid) has an account at branch (bid):
 $\pi_{cid,bid} (\text{customer} \triangleright \triangleleft \text{account})$
 - Divide the above by **bid**'s of all branches in Westmount
 $\pi_{bid} (\sigma_{district = "Westmount"} (\text{branch}))$
- That is: $\pi_{name} ((\text{customer} \triangleright \triangleleft \text{account}) \div \pi_{bid} (\sigma_{district = "Westmount"} (\text{branch})))$
- The division “**r** \div **s**” of relation r by s is defined as:
$$r \div s = \pi_{R-S} (\mathbf{r}) - \pi_{R-S} (\pi_{R-S} (\mathbf{r}) \times \mathbf{s} - \mathbf{r})$$