

# COMP353 Databases

---

**More on SQL:**

**Nested Querye  
Views**

# Scalar Values

---

- An SQL query is an expression that evaluates to a **collection** of tuples, i.e., it produces a **relation/bag**
- This “collection” may have only one attribute
- It is also possible that there will be only one single value produced for that attribute
- If all these hold, then we say that the query produces a **scalar** value
  - **Scalar values** – example include simple values such as integers, reals, strings, dates, etc.

# Queries that Produce Scalar Values

---

- Relation schema:

**Movie**( title, year, length, filmType, studioName, producerC# )

- Query:

Find **certificate number** of the producer of “**Star Wars**”

- Query in SQL:

```
SELECT producerC#  
FROM Movie  
WHERE title = 'Star Wars';
```

*Assuming that we have only one such movie.*

# Subqueries

---

- Conditions in the **WHERE** clause may have comparisons that involve scalar values
  - A SQL query can produce a scalar value
  - If so, we can use such SELECT-FROM-WHERE expression, surrounded by parentheses, **as if it were a constant**
  - **Subquery** – a query within a query
- The result of a SQL subquery is a collection (relation/bag)

# Example

---

- Relation schemas:

**Movie** (title, year, length, filmType, studioName, producerC#)

**Exec** (name, address, cert#, netWorth)

- Query:

Find the name of the producer of “Star Wars”

- Query in SQL:

**SELECT** Exec.name

**FROM** Movie, Exec

**WHERE** Movie.title = 'Star Wars' **AND**

Movie.producerC# = Exec.cert#;

# Example

---

- Relation schemas:

**Movie** (title, year, length, filmType, studioName, producerC#)

**Exec** (name, address, cert#, netWorth)

- Query:

Find the name of the producer of “Star Wars”

- Query with Subquery:

**SELECT** name

**FROM** Exec

**WHERE** cert# = ( **SELECT** producerC#  
                  **FROM** Movie  
                  **WHERE** title = 'Star Wars' );

# Conditions Involving Relations

---

- There are a number of SQL checks/conditions that can be done on a relation **R** and produce a **boolean** value
- These conditions can be **negated** by putting a **NOT** before them
- Typically, **R** above is the result of an SQL subquery, shown as: **(R)**
- If such a condition involves a scalar value **s** or a tuple, we should make sure its type matches **R**.

# Conditions Involving Relations

- “**EXISTS ( $R$ )**” is a condition that is true iff  $R$  is not empty
- “ **$s$  IN ( $R$ )**” is true iff  $s$  is equal to **one** of the values in  $R$ 
  - “ **$s$  NOT IN ( $R$ )**” is true iff  $s$  is not equal to any value in  $R$
- “ **$s$  > ALL ( $R$ )**” is true iff  $s$  is greater than **every** value in  $R$ 
  - “>” could be replaced by other operators with the analogous meaning
  - **Note:** “ **$s$  <> ALL ( $R$ )**” is the same as “ **$s$  NOT IN  $R$** ”
- “ **$s$  > ANY ( $R$ )**” is true iff  $s$  is > **at least one** value in  $R$ 
  - “>” could be replaced by any of the other 5 comparison operators with the analogous meaning
  - **Note:** “ **$s$  = ANY ( $R$ )**” is the same as “ **$s$  IN  $R$** ”



# Conditions Involving Tuples

- A tuple in SQL is represented by a parenthesized list of scalar values; the concept “tuple” can be viewed as an *extension* of the concept of scalar;
  - (123, 'foo')
- Mixing of constants and attributes is also permitted in tuples
  - (123, Movie.title)
- If a tuple  $t$  has the same number of components as a relation  $R$ , then it makes sense to compare  $t$  and  $R$  like:
  - $t \text{ IN } (R)$  -- this is true iff  $t$  is in  $R$
  - $t \neq \text{ANY } (R)$  -- this is true  $R$  includes a tuple other than  $t$

# Example

- Relation schemas:

**Movie** (title, year, length, filmType, studioName, producerC#)

**Exec** (name, address, cert#, netWorth)

**StarsIn** (title, year, starName)

- Query: Find the names of the producers of Harrison Ford's movies

- Query in SQL:

**SELECT** name

**FROM** Exec

**WHERE** cert# **IN** (**SELECT** producerC#

**FROM** Movie

**WHERE** (title, year) **IN** (**SELECT** title, year

**FROM** StarsIn

**WHERE** starName = 'Harrison Ford'));

# Example

- Relation schemas:  
**Movie**(title, year, length, filmType, studioName, producerC#)  
**Exec**(name, address, cert#, netWorth)  
**StarsIn**(title, year, starName)
- Query: Find names of the producers of Harrison Ford's movies
- Query in SQL:  
**SELECT** Exec.name  
**FROM** Exec, Movie, StarsIn  
**WHERE** Exec.cert# = Movie.producerC# **AND**  
          Movie.title = StarsIn.title **AND**  
          Movie.year = StarsIn.year **AND**  
          starName = 'Harrison Ford';

# Correlated Subqueries

---

- Simple subqueries can be evaluated once and the result be used in a higher level (calling) query
- A more complex use of nested subquery requires the subquery to be evaluated many times, once for each assignment of a value (to some term in the subquery) that comes from a tuple variable in the calling query
- A subquery of this type is called ***correlated subquery***

# Correlated Subqueries

- Relation schema:  
**Movie**(title, year, length, filmType, studioName, producerC#)
- Query:  
Find movie titles that appear more than once
- Query in SQL:  
**SELECT** title  
**FROM** Movie Old  
**WHERE** year < ANY (**SELECT** year  
                            **FROM** Movie  
                            **WHERE** title = Old.title);  
Note the **scopes** of the variables in this query.

# Correlated Subqueries

- ## ■ Query in SQL

# SELECT title

## FROM Movie Old

```
WHERE year < ANY (SELECT year
                  FROM Movie
                  WHERE title = Old.title);
```

- The condition in the outer **WHERE** is true only if there is a movie with same title as *Old.title* that has a **later** year
  - ➔ The query will produce a title **one fewer times** than there are movies with that title
- What would be the result if we used “<>”, instead of “<” ?
  - ➔ For a movie title appearing 3 times, we would get 3 copies of the title in the output

# Views

- **View** is a table/relation defined in a database but has no tuples explicitly stored for it in the database but rather computed, when needed, from the **view definition**
- The view mechanism provides support for:
  - Logical data independence:
    - Views can be used to define relations in the external schema that mask, from the applications/users, changes in the *conceptual database schema*
    - If the schema of a relation is changed, we can define a view with the old schema so that applications that use the old schema can continue using it
  - Security:
    - Views can be used to restricts the users access only the information they are allowed to “see and operate on”

# Views

---

- Relation schema:

**Movie**(title, year, length, filmType, **studioName**, producerC#)

- View:

Create the **Paramount's movies** (title and year)

- View in SQL:

**CREATE VIEW** ParamountMovie **AS**

**SELECT** title, year

**FROM** Movie

**WHERE** studioName = 'Paramount';



# Views

---

- A view can be used in defining new queries/views in exactly the same way as an explicitly stored table may be used
- Example to query the (virtual) relation ParamountMovie

```
SELECT title  
FROM ParamountMovie  
WHERE year = 1979;
```

- This query is translated, by the query processor, into:

```
SELECT title  
FROM Movie  
WHERE studioName = 'Paramount' AND year = 1979;
```

# Views

---

- Relation schema:

ParamountMovie (title, year )

StarsIn(title, year, starName)

- Query:

List the stars of the movies made by Paramount.

- Query in SQL

**SELECT DISTINCT** StarsIn.starName

**FROM** ParamountMovie, StarsIn

**WHERE** ParamountMovie.title = StarsIn.title **AND**

ParamountMovie.year = StarsIn.year;

# Views

---

- Relation schema:

**Movie** (title, year, length, filmType, studioName, producerC#)

**Exec** (name, address, cert#, netWorth)

- View:

Define a view of Movie (titles and executives/producers)

- View in SQL:

**CREATE VIEW** MovieProd **AS**

**SELECT** Movie.title, Exec.name

**FROM** Movie, Exec

**WHERE** Movie.producerC# = Exec.cert#;

# Views

---

- Relation schema:

**MovieProd** (title, name)

- Query:

Find the name of the producer of 'Gone With the Wind'?

- Query in SQL:

**SELECT** name

**FROM** MovieProd

**WHERE** title = 'Gone With the Wind';

# Views

---

- Renaming attributes used in view definitions
  - We can give new names to view attributes rather than using the names that come out of query defining the view

- Example:

```
CREATE VIEW MovieProd (MovieTitle, ProducerName) AS  
SELECT Movie.title, Exec.name  
FROM Movie, Exec  
WHERE Movie.producerC# = Exec.cert#;
```

# Views

---

- Relation schema:

MovieProd (MovieTitle, ProducerName)

- Query:

Find the name of the producer of 'Gone With the Wind'?

- Query in SQL:

**SELECT** ProducerName

**FROM** MovieProd

**WHERE** MovieTitle = 'Gone With the Wind';

# Updating Views?

---

- We saw that a view can appear in a query in exactly the same way as a “base” table may appear.
- What about modifications/updates?
- What does it mean to update a view?
  - Translate modification of the view to the corresponding modification on the base tables used in the view definition
- Should we allow updates on views?
  - Yes, in principle, but some problems may arise
- Some “simple” views can be updated
  - Such views are called **updatable views**
- Many views **cannot** be updated
  - This is due to the so called **view-update anomaly**

# Insertion into Views?

- Relation schema:

**Movie**(title, year, length, filmType, **studioName**, producerC#)

- View: Recall the definition of **ParamountMovie**

**CREATE VIEW** ParamountMovie **AS**

**SELECT** title, year

**FROM** Movie

**WHERE** studioName = 'Paramount';

- Update statement:

**INSERT INTO** ParamountMovie (**title,year**) **VALUES**('KK', 2002);

- Result: the following tuple being added to Movie

('KK', 2002, NULL, NULL, **NULL**, NULL) **What's the problem?**



# Insertion into Views?

- Relation schema:

**Movie**(title, year, length, filmType, studioName, producerC#)

- An updatable view:

**CREATE VIEW** ParamountMovie **AS**

**SELECT** title, year, studioName

**FROM** Movie

**WHERE** studioName = 'Paramount';

- Update statement:

**INSERT INTO** ParamountMovie **VALUES**('KK',2002,'Paramount');

- Result: the following tuple is being added to Movie

('KK', 2002, NULL, NULL, 'Paramount', NULL) Problem solved!

# Insertion into Views?

- Relation schemas:

**Movie**(title, year, length, filmType, studioName, producerC# )

**Exec**(name, address, cert#, netWorth)

- View in SQL:

**CREATE VIEW** MovieProd **AS**

**SELECT** Movie.title, Exec.name

**FROM** Movie, Exec

**WHERE** Movie.producerC# = Exec.cert#;

- Update statement

**INSERT INTO** MovieProd (**title,name**) **VALUES**('The Movie', 'J. Smith');

- Result: these tuples are added to the corresponding relations:

**Movie**('The Movie', **NULL**, NULL, NULL, NULL, **NULL**)

**Exec**('J. Smith', NULL, **NULL**, NULL)    **Problems? The insertion command will fail !**

# Deletion from Views?

- Relation schema:

**Movie**(title, year, length, filmType, studioName, producerC#)

- View: Recall the definition :

**CREATE VIEW** ParamountMovie **AS**

**SELECT** title, year, studioName

**FROM** Movie

**WHERE** studioName = 'Paramount';

- Delete statement:

**DELETE FROM** ParamountMovie **WHERE** title **LIKE** '%K%';

- Translated query:

**DELETE FROM** Movie

**WHERE** studioName = 'Paramount' **AND** title **LIKE** '%K%';

# Updating Views?

---

- Relation schema:

**Movie**(title, year, length, filmType, studioName, producerC#)

- View:

**CREATE VIEW** ParamountMovie **AS**

**SELECT** title, year, studioName

**FROM** Movie

**WHERE** studioName = 'Paramount';

- The view update statement:

**UPDATE** ParamountMovie **SET** year = 1797 **WHERE** title = 'KK';

- We may drop a view: **DROP VIEW** ParamountMovie;

# Updating Views?

---

- Recall: updating views includes insertion, deletion, and changing data
- SQL provides a formal definition of when modifications to a view are permitted
- Roughly, this is permitted when the view is defined by selecting some attributes from **one** relation **R**, which could be an “updatable” view itself
  - The list in the **SELECT** clause includes “enough” attributes that for every tuple inserted into the view, the tuple inserted into the base relation will “yield” the inserted tuple of the view
  - The **NOT NULL** constraints on the base table will not be violated
  - The view definition uses **SELECT** (but not **SELECT DISTINCT**)
  - The **WHERE** clause does not involve **R** in a subquery